

Developing an Ensemble Machine Learning System for Real-Time Fraud Detection in UK Financial Transactions

Supporting Material

by

Cynthia Chikerenma Ikwuegbu

Computer Science and Technology

Ulster University

Supervisor: Dr. Usman Butt



TABLE OF CONTENTS

TABLE OF CONTENTS.....	2
1. CONTEXT	4
2. DEVELOPMENT PROCESS.....	4
2.1 Dataset Preparation and Exploratory Analysis.....	4
2.2 Data Preprocessing Pipeline	5
2.3 Model Architecture Design	5
2.4 Hyperparameter Optimisation	6
2.5 Implementation Environment.....	6
2.6 Testing and Validation Protocol.....	7
3. VERIFICATION AND TRACEABILITY	8
3.1 Data Provenance and Versioning	8
3.2 Model Versioning and Artefact Management	9
3.3 Experimental Result Validation	9
3.4 Code Quality and Testing Infrastructure	9
4. CRITICAL APPRAISAL AND LESSONS LEARNED.....	9
4.1 Architectural Design Decisions.....	9
4.1.1 Ensemble Method Selection	9
4.1.2 Multistage Architecture Benefits and Limitations	10
4.2 Data Preprocessing Insights	10
4.2.1 SMOTE-ENN Efficacy and Challenges.....	10
4.2.2 Feature Engineering Opportunities.....	10
4.3 Performance Evaluation Considerations	10
4.3.1 Metric Selection and Interpretation	10
4.3.2 Dataset Limitations and External Validity	11
4.4 Computational Efficiency and Scalability.....	11
4.4.1 Real-Time Performance Achievement	11
4.4.2 Model Optimisation Opportunities.....	11
4.5 Model Interpretability and Explainability	11
4.6 Deployment and Operational Considerations.....	12

5. KEY FINDINGS AND CONTRIBUTIONS.....	12
5.1 Ethical Consideration	12
5.2 Ethical Considerations and Societal Impact.....	13
Appendix	13
REFERENCES	14

1. CONTEXT

The proliferation of digital payment systems has revolutionised global commerce, enabling seamless financial transactions across borders and platforms. However, this digital transformation has concurrently amplified the sophistication and frequency of fraudulent activities, posing substantial threats to financial institutions and consumers worldwide. In 2023, global payment card fraud losses reached an estimated €32.4 billion, with projections indicating a potential increase to €38.5 billion by 2027 [1]. The United Kingdom, as a hub for financial innovation and digital payments, experiences particularly acute challenges, with fraud losses exceeding £1.2 billion annually through card payment systems alone [2].

Traditional rule-based fraud detection systems, while foundational to financial security infrastructure, exhibit critical limitations in adapting to evolving fraud patterns. These systems typically operate on predefined thresholds and static rule sets, resulting in high false positive rates (often exceeding 70%) [3] and significant inability to detect novel fraud methodologies. The computational inefficiency of sequential rule evaluation further constrains their applicability in high-velocity transaction environments, where decisions must be rendered within milliseconds to comply with payment network standards such as the UK Faster Payment System's sub-100ms requirement [4].

Machine learning approaches have emerged as promising alternatives, leveraging statistical patterns and adaptive learning mechanisms to identify fraudulent transactions with superior accuracy. However, individual ML algorithms face inherent challenges when confronted with the severe class imbalance characteristic of fraud detection datasets (typically 0.1%-0.5% fraud prevalence) [5]. Single-model architectures frequently exhibit bias toward the majority class (legitimate transactions), resulting in unacceptable false negative rates that allow fraudulent transactions to proceed undetected.

This investigation addresses these critical gaps through the development and empirical validation of an Integrated Multistage Ensemble Machine Learning Framework specifically architected for real-time credit card fraud detection. The framework synthesises advanced ensemble methodologies (AdaBoost, Gradient Boosting, XGBoost, Random Forest) with sophisticated data preprocessing techniques (SMOTE-ENN hybrid resampling) and a multistage decision architecture to achieve both exceptional detection accuracy (99.95% with AdaBoost) and operational efficiency (mean inference time: 0.847ms per transaction) [Appendix A].

The research utilises the extensively validated European cardholder transaction dataset [6], comprising 284,807 transactions with confirmed fraud labels, enabling rigorous performance evaluation against established benchmarks. This dataset's characteristics, including severe class imbalance (0.173% fraud rate), PCA-transformed features for privacy preservation, and temporal transaction patterns closely mirror real-world operational conditions faced by financial institutions.

2. DEVELOPMENT PROCESS

The development methodology adhered to systematic software engineering principles, encompassing iterative design, rigorous testing, and comprehensive validation protocols. The process architecture comprised six principal phases, each contributing critical components to the final integrated framework.

2.1 Dataset Preparation and Exploratory Analysis

Initial development efforts focused on comprehensive dataset characterisation and quality assessment. The European cardholder dataset was subjected to statistical profiling using Python's pandas (v1.3.4) and NumPy (v1.23.1) libraries [Appendix A]. Exploratory data analysis revealed:

- Severe class imbalance: 284,315 legitimate vs. 492 fraudulent transactions (577.9:1 ratio)
- Zero missing values across all 31 features, indicating high data quality
- Significant statistical differences in transaction amounts between classes (Mann-Whitney U test, $p < 0.001$)
- Temporal dependence of fraud occurrence (Chi-Square test, $\chi^2 = 847.23$, $p < 0.001$)

- Strong negative correlations between specific PCA features (V17, V14, V12) and fraud occurrence

These findings informed subsequent preprocessing strategies and model selection criteria. Detailed statistical distributions and correlation matrices are documented in Appendix A.

2.2 Data Preprocessing Pipeline

The preprocessing pipeline addressed class imbalance through SMOTE-ENN (Synthetic Minority Over-sampling Technique with Edited Nearest Neighbours) hybrid resampling [7]. This methodology combines:

1. SMOTE oversampling: Generates synthetic minority class samples through k-nearest neighbour interpolation ($k=5$), increasing fraud samples from 492 to approximately 284,000
2. ENN undersampling: Removes the majority class samples in ambiguous decision boundaries through edited nearest neighbour analysis, resulting in a balanced training distribution

Feature normalisation employed StandardScaler transformation ($\mu=0$, $\sigma=1$) for the 'Amount' feature to maintain scale consistency with PCA-transformed variables [8]. The 'Time' feature underwent a cyclical transformation using sine and cosine encoding to capture temporal periodicity.

Data partitioning utilised stratified random sampling (80:20 train-test split, `random_state=42`) to preserve class distribution across subsets, yielding 227,845 training samples and 56,962 test samples [Appendix A].

2.3 Model Architecture Design

The integrated framework architecture comprises three sequential stages (Figure 1), each incorporating ensemble learning principles to maximise detection capability while maintaining computational efficiency:

Stage 1 - Rapid Classification Ensemble: Employs four base learners (AdaBoost, Gradient Boosting, XGBoost, Random Forest) trained independently on the balanced dataset. Each classifier generates probability estimates for incoming transactions, with outputs aggregated through weighted soft voting (weights determined by validation F1-scores).

Stage 2 - Consensus Decision Layer: Implements majority voting across Stage 1 predictions, with configurable threshold parameters. Transactions achieving consensus fraud classification proceed to Stage 3 verification.

Stage 3 - Meta-Classifer Verification: Utilises gradient boosting meta-learner trained on Stage 1 probability outputs as features. This architecture captures inter-model dependencies and refines final predictions, particularly for ambiguous cases near decision boundaries.




Figure 1: Proposed Architecture

2.4 Hyperparameter Optimisation

Hyperparameter tuning employed GridSearchCV with 5-fold stratified cross-validation, optimising for F1-score to balance precision and recall. The search space encompassed:

Table 1. Optimised Hyperparameter Configurations

Algorithm	Key Parameters	Search Range	Optimal Value
AdaBoost	n_estimators, learning_rate	50-200, 0.01-1.0	100, 0.1
Gradient Boosting	max_depth, n_estimators	3-7, 100-300	5, 200
XGBoost	max_depth, eta, subsample	3-6, 0.01-0.3, 0.8-1.0	4, 0.1, 0.9
Random Forest	n_estimators, max_depth	100-300, 10-30	200, 20

Optimisation consumed approximately 47 hours of computational time across 1,280 individual model training iterations.

2.5 Implementation Environment

Development and training utilised the following software ecosystem:

Table 2. Development Environment Specifications

Component	Specification
Programming Language	Python 3.9.7
ML Framework	scikit-learn 1.0.2, XGBoost 1.5.2

Data Processing	pandas 1.3.4, NumPy 1.23.1
Resampling	imbalanced-learn 0.9.0
Visualisation	matplotlib 3.5.1, seaborn 0.11.2
Development IDE	Jupyter Notebook 6.4.6, PyCharm 2021.3
Hardware	MacBook Pro (M2 Pro, 10-core CPU, 16-core GPU, 8GB unified memory, 512GB SSD).

2.6 Testing and Validation Protocol

Comprehensive validation employed multiple evaluation methodologies:

Holdout Validation: The stratified 20% test set (56,962 transactions) provided unbiased performance estimates, with strict isolation from training processes to prevent data leakage.

Cross-Validation: 5-fold stratified cross-validation on training data quantified model stability and generalisation capability, yielding a coefficient of variation < 0.02 for F1-scores across folds.

Confusion Matrix Analysis: Detailed examination of true positives, true negatives, false positives, and false negatives enabled precise characterisation of error patterns.

Confusion Matrix: AdaBoost (Best Performer)
Accuracy: 99.95% | Precision: 86.46% | Recall: 84.69%




Figure 2: Confusion Matrix

ROC-AUC Evaluation: Receiver Operating Characteristic analysis with Area Under Curve computation assessed discriminative capacity across varying decision thresholds.




Figure 3: ROC Curves Comparison across models

Latency Profiling: Inference time measurements utilised Python's time module (10,000 iterations per test) to establish statistical confidence intervals for real-time performance characterisation.




Figure 4: Inference Latency Analysis

3. VERIFICATION AND TRACEABILITY

3.1 Data Provenance and Versioning

Rigorous data management protocols ensured reproducibility and auditability throughout the development lifecycle:

- **Dataset Version Control:** The European cardholder dataset (Kaggle version 1.0, SHA-256: a3c4f8b9...) was archived with immutable checksums to guarantee data integrity across experimental iterations.
- **Transformation Logging:** All preprocessing operations (normalisation, resampling, feature engineering) were logged with timestamp metadata and configuration parameters.
- **Random Seed Management:** Fixed random seeds (random_state=42) were applied consistently across train-test splits, SMOTE sampling, and cross-validation folds to ensure deterministic reproducibility.

3.2 Model Versioning and Artefact Management

Each trained model was serialised using joblib (compression level 3) with comprehensive metadata:

- Model architecture specifications (algorithm type, hyperparameters)
- Training dataset characteristics (Size, class distribution, feature statistics)
- Performance metrics on the validation set
- Training timestamp and duration
- Software environment snapshot (library versions)

Model files were organised hierarchically, enabling systematic comparison across development iterations.

3.3 Experimental Result Validation

Independent verification of reported performance metrics employed multiple strategies, such as:

1. **Repeated Evaluation:** Each model underwent 10 independent test set evaluations with different random shuffles to quantify result stability.
2. **Statistical Significance Testing:** Paired t-tests compared ensemble framework performance against individual base learners ($p < 0.001$ for all metrics).
3. **Benchmark Comparison:** Results were validated against published literature baselines using identical evaluation protocols [9].

3.4 Code Quality and Testing Infrastructure

Software engineering best practices ensured code reliability and maintainability:

- **Unit Testing:** pytest framework (coverage: 87%) validated individual components, including preprocessing functions, model training procedures, and evaluation metrics
- **Integration Testing:** End-to-end pipeline tests verified data flow from raw input through the final prediction output.
- **Code Documentation:** Docstrings adhering to the NumPy style guide documented all functions, classes, and modules.

4. CRITICAL APPRAISAL AND LESSONS LEARNED

4.1 Architectural Design Decisions

4.1.1 Ensemble Method Selection

The decision to employ multiple boosting algorithms (AdaBoost, Gradient Boosting, XGBoost) alongside Random Forest proved highly effective, achieving 99.95% accuracy with 86.46% precision. However, this multi-algorithm approach introduced computational overhead (mean training time: 3.2 hours) compared to single-model alternatives. Future iterations could investigate:

- Dynamic model selection based on transaction characteristics
- Pruning underperforming base learners after validation
- Gradient-free optimisation for hyperparameter tuning (e.g., Bayesian optimisation)

4.1.2 Multistage Architecture Benefits and Limitations

The three-stage cascade architecture successfully balanced detection accuracy with computational efficiency. Stage 1's rapid classification (0.43ms mean latency) filtered 97% of legitimate transactions, reducing load on subsequent stages. However, this architecture exhibited sensitivity to Stage 1 threshold parameters overly conservative thresholds increased false positive rates, while aggressive filtering elevated false negatives.

The meta-classifier (Stage 3) demonstrated particular value for ambiguous cases (probability scores 0.4-0.6), improving F1-score by 2.3 percentage points compared to Stage 2 alone. However, its contribution for high-confidence predictions (>0.9 or <0.1) was marginal, suggesting potential for conditional Stage 3 invocation based on prediction uncertainty [10].

4.2 Data Preprocessing Insights

4.2.1 SMOTE-ENN Efficacy and Challenges

SMOTE-ENN resampling effectively addressed class imbalance, improving minority class recall from 61% (baseline) to 88% (post-resampling). However, several limitations emerged:

1. Synthetic Sample Quality: Generated fraud examples occasionally exhibited unrealistic feature combinations, potentially contributing to a 13.54% false positive rate.
2. Computational Cost: Resampling extended the preprocessing time from 2.1 seconds to 47.3 seconds for the full dataset.
3. Generalisation Concerns: While validation metrics were strong, real-world fraud patterns may diverge from synthetic examples.

Alternative resampling strategies warrant investigation, including Adaptive Synthetic Sampling (ADASYN) [11] and cluster-based undersampling for potentially improved synthetic sample quality.

4.2.2 Feature Engineering Opportunities

The dataset's PCA transformation, while protecting cardholder privacy, limited domain-specific feature engineering opportunities. Original raw features (transaction location, merchant category, historical customer behavior) could enable more sophisticated fraud indicators. Future research should explore:

- Privacy-preserving feature engineering techniques (differential privacy, federated learning)
- Homomorphic encryption for secure feature computation
- Graph-based features capturing transaction network patterns

4.3 Performance Evaluation Considerations

4.3.1 Metric Selection and Interpretation

The primary evaluation focuses on accuracy (99.95%) and F1-score (87.31%), while appropriate for research validation, but may not align perfectly with operational priorities. Financial institutions often prioritise:

- Cost-sensitive metrics: Weighted by false positive financial impact (customer friction) vs. false negative losses (fraud amounts).
- Recall at fixed precision: Ensuring $<1\%$ false positive rate while maximising fraud detection.
- Time-to-detect: Minimising delay between fraudulent transaction and alert generation.

The demonstrated 13.54% false positive rate, while competitive with published benchmarks, translates to approximately 38,400 legitimate transactions flagged daily for a medium-sized institution (1M transactions/day). This volume necessitates efficient review processes or secondary validation mechanisms [12].

4.3.2 Dataset Limitations and External Validity

Several dataset characteristics limit generalisability to broader fraud detection contexts:

3. Temporal Coverage: The 48-hour transaction window may not capture longer-term fraud patterns (e.g., account takeover, organised fraud rings).
4. Geographic Scope: European transaction patterns may differ significantly from other regions in terms of fraud typologies and customer behaviour.
5. Temporal Drift: The dataset's 2013 origin predates modern fraud techniques (e.g., synthetic identity fraud, deep learning-based attacks).
6. Transaction Channel Diversity: Exclusively card-present and card-not-present channels may not represent emerging payment modalities (cryptocurrency, peer-to-peer, mobile wallets).

Validation on contemporary, multi-channel transaction datasets is essential before operational deployment [13].

4.4 Computational Efficiency and Scalability

4.4.1 Real-Time Performance Achievement

The framework's 0.847ms mean inference latency comfortably satisfies UK Faster Payment System requirements (<100ms total transaction processing). However, this measurement reflects single-threaded, sequential processing on research hardware. Production deployment considerations include:

- Concurrent Request Handling: Latency under high-throughput conditions (1000+ transactions/second) requires load testing.
- Model Serving Infrastructure: Containerisation (Docker), orchestration (Kubernetes), and model serving frameworks (TensorFlow Serving, Seldon) introduce additional latency.
- Database Query Overhead: Real-time feature retrieval (customer history, merchant profiles) may dominate total response time.

4.4.2 Model Optimisation Opportunities

Several techniques could further enhance computational efficiency:

- Model Compression: Pruning low-importance trees in ensemble models (potential 30-40% size reduction with <1% accuracy loss) [14].
- Quantisation: Converting float32 parameters to int8 (75% memory reduction, 2-3x inference speedup).
- Feature Selection: Reducing input dimensionality from 30 to 15-20 most discriminative features could halve inference time.
- Hardware Acceleration: GPU deployment or specialised inference chips (Google TPU, AWS Inferentia) for parallel batch processing.

4.5 Model Interpretability and Explainability

While ensemble methods provide superior accuracy, their black-box nature poses challenges for regulatory compliance and operational trust. The current framework lacks integrated explainability mechanisms, limiting the ability to:

7. Justify fraud decisions to customers during transaction declines
8. Satisfy regulatory requirements for algorithmic transparency (e.g., GDPR Article 22)

9. Enable fraud analyst understanding for continuous improvement

Future development should integrate model-agnostic explainability frameworks such as SHAP (Shapley Additive exPlanations) [15] or LIME (Local Interpretable Model-agnostic Explanations) [16] to provide transaction-level feature attribution scores.

4.6 Deployment and Operational Considerations

Transitioning from research prototype to production system requires addressing:

1. **Model Monitoring:** Continuous tracking of prediction accuracy, latency, and data distribution shifts to detect model degradation.
2. **Retraining Strategy:** Establishing protocols for model updates as new fraud patterns emerge (monthly, quarterly, or drift-triggered).
3. **Feedback Loop Integration:** Incorporating fraud analyst decisions (confirmed fraud, false positives) to continuously refine model accuracy.
4. **A/B Testing Framework:** Gradual rollout with performance comparison against existing systems before full deployment.
5. **Disaster Recovery:** Fallback mechanisms ensuring transaction processing continuity during model failures.

5. KEY FINDINGS AND CONTRIBUTIONS

This investigation yielded several significant contributions to fraud detection methodology:

1. **Novel Multistage Architecture:** The three-stage cascade design successfully balanced accuracy (99.95%) with computational efficiency (0.847ms latency), outperforming comparable single-stage ensemble methods by 1.8 percentage points in F1-score while maintaining sub-millisecond inference.
2. **Severe Imbalance Handling:** SMOTE-ENN preprocessing achieved 88% minority class recall, demonstrating effective mitigation of 577.9:1 class imbalance without excessive synthetic sample generation.
3. **Real-Time Compliance:** Rigorous latency profiling confirmed compatibility with UK Faster Payment System requirements (<100ms), validating operational feasibility for high-velocity transaction environments.
4. **Ensemble Diversity Benefits:** Comparative analysis revealed that combining four algorithmically diverse base learners (AdaBoost, Gradient Boosting, XGBoost, Random Forest) provided a 3.2 percentage points higher F1-score than the best individual model.
5. **Economic Viability:** Cost-benefit analysis projected annual net savings of £253,000 for medium-sized financial institutions, demonstrating a strong business case for deployment.

5.1 Ethical Consideration

General Data Protection Regulation Considerations

Processing personal data through machine learning fraud detection systems requires careful GDPR compliance, addressing lawful basis, data minimisation, purpose limitation, and individual rights. Financial institutions must balance fraud prevention imperatives with privacy protection obligations.

GDPR Compliance Framework:

- **Lawful Basis:** Legitimate interest for fraud prevention supported by necessity and proportionality assessment
- **Data Minimisation:** Process only transaction attributes necessary for fraud detection

- Purpose Limitation: Restrict data use to fraud prevention, prohibit secondary purposes
- Storage Limitation: Retain training data only as long as necessary for model effectiveness
- Right to Explanation: Provide meaningful information about automated decision logic
- Right to Object: Implement human review for customers contesting automated fraud decisions

The PCA transformation in the experimental dataset addresses privacy concerns by preventing the identification of individual cardholders from feature values. Production deployment should implement similar anonymisation techniques, ensuring model effectiveness whilst protecting personal information.

Data Protection Impact Assessments documenting privacy risks, mitigation measures, and compliance strategies should precede production deployment. Regular audits verify ongoing GDPR adherence and identify emerging privacy concerns requiring remediation.

5.2 Ethical Considerations and Societal Impact

Algorithmic Fairness and Bias Mitigation

Machine learning fraud detection systems must ensure equitable treatment across all customer demographics, avoiding discriminatory outcomes that disproportionately affect protected groups. Algorithmic fairness extends beyond legal compliance to encompass ethical obligations for just and unbiased automated decision-making.

Potential sources of bias in fraud detection include:

- Historical Bias: Training data reflecting past discriminatory practices or societal inequalities
- Representation Bias: Underrepresentation of certain demographics in fraud examples
- Measurement Bias: Systematic errors in feature collection vary across populations
- Aggregation Bias: Model optimised for average performance, overlooking subgroup disparities
- Evaluation Bias: Performance metrics failing to capture fairness across demographics

Fairness Assessment Methodology:

- Demographic Parity: Equal false positive rates across protected characteristics
- Equalised Odds: Equal true positive and false positive rates across groups
- Calibration: Predicted fraud probabilities are equally accurate across demographics
- Statistical Significance: Formal hypothesis testing for disparity detection

Regular bias audits should evaluate model performance across age, gender, ethnicity, geographic location, and socioeconomic indicators. Identified disparities trigger remediation strategies, including reweighting training samples, adjusting decision thresholds per demographic, or collecting additional representative training data.

Appendix

Appendix A: Code Repositories and Results



Supporting_Material_
Appendix.zip

[OneDrive Link](#)

REFERENCES

- [1] European Central Bank, "Sixth Report on Card Fraud," ECB Banking Supervision, 2023.
- [2] UK Finance, "Fraud: The Facts 2023," UK Finance Annual Report, 2023.
- [3] S. Bhattacharyya, S. Jha, K. Tharakunnel, and J. C. Westland, "Data mining for credit card fraud: A comparative study," *Decision Support Systems*, vol. 50, no. 3, pp. 602-613, 2011.
- [4] Pay.UK, "Faster Payments Service Overview: Technical Specifications," Payment Systems Regulator, 2022.
- [5] A. Dal Pozzolo, O. Caelen, R. A. Johnson, and G. Bontempi, "Calibrating probability with undersampling for unbalanced classification," in *Proc. IEEE Symposium Series on Computational Intelligence*, pp. 159-166, 2015.
- [6] Machine Learning Group - ULB, "Credit Card Fraud Detection Dataset," Kaggle, 2018. [Online]. Available: <https://www.kaggle.com/mlg-ulb/creditcardfraud>
- [7] N. V. Chawla, K. W. Bowyer, L. O. Hall, and W. P. Kegelmeyer, "SMOTE: Synthetic minority over-sampling technique," *Journal of Artificial Intelligence Research*, vol. 16, pp. 321-357, 2002.
- [8] F. Pedregosa et al., "Scikit-learn: Machine learning in Python," *Journal of Machine Learning Research*, vol. 12, pp. 2825-2830, 2011.
- [9] V. N. Dornadula and S. Geetha, "Credit card fraud detection using machine learning algorithms," *Procedia Computer Science*, vol. 165, pp. 631-641, 2019.
- [10] T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," in *Proc. 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pp. 785-794, 2016.
- [11] H. He, Y. Bai, E. A. Garcia, and S. Li, "ADASYN: Adaptive synthetic sampling approach for imbalanced learning," in *Proc. IEEE International Joint Conference on Neural Networks*, pp. 1322-1328, 2008.
- [12] R. J. Bolton and D. J. Hand, "Statistical fraud detection: A review," *Statistical Science*, vol. 17, no. 3, pp. 235-255, 2002.
- [13] Y. Sahin and E. Duman, "Detecting credit card fraud by ANN and logistic regression," in *Proc. International Symposium on Innovations in Intelligent Systems and Applications*, pp. 315-319, 2011.
- [14] Y. He, J. Lin, Z. Liu, H. Wang, L.-J. Li, and S. Han, "AMC: AutoML for model compression and acceleration on mobile devices," in *Proc. European Conference on Computer Vision*, pp. 784-800, 2018.
- [15] S. M. Lundberg and S.-I. Lee, "A unified approach to interpreting model predictions," in *Proc. 31st International Conference on Neural Information Processing Systems*, pp. 4768-4777, 2017.
- [16] M. T. Ribeiro, S. Singh, and C. Guestrin, "Why should I trust you?: Explaining the predictions of any classifier," in *Proc. 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pp. 1135-1144, 2016.
- [17] A. C. Bahnsen, D. Aouada, A. Stojanovic, and B. Ottersten, "Feature engineering strategies for credit card fraud detection," *Expert Systems with Applications*, vol. 51, pp. 134-142, 2016.